We vs Them

Detecting Othering Language in Social Media — Mid-Internship Progress · May 2026

"Othering" is the discursive construction of an in-group ("we") against a dehumanised out-group ("them"). The goal of this project is to build an end-to-end NLP pipeline that detects such language at scale in social media posts, combining rule-based linguistics and machine learning.

134 881	**10 000**	**7 115**	**5.3 %**	**Step 4**
posts collected	posts ML-scored	othering posts flagged	othering rate	classifier — in progress

Step 0 — Data Collection

■ **Source 1** — ucberkeley-dlab/measuring-hate-speech (HuggingFace)

→ 133 592 annotated social media comments

→ Multi-label hate-speech ratings by trained annotators

■ **Source 2** — Pushshift Reddit (streaming, 50 k rows)

→ 1 289 posts after subreddit filtering

→ Subreddits: *politics, immigration, europe, worldnews, conspiracy*

■ **Cleaning**

→ Title + body merged for Reddit posts

→ Posts < 20 chars removed

→ Unified schema: text · source · subreddit

■ **Final dataset**

→ **134 881 posts** saved to *dataset_clean.csv*

Steps 1 & 2 — Text Enrichment

■ **Pronoun tagging** (all 134 k posts)

→ *We-markers*: we, us, our, ourselves

→ *Them-markers*: they, them, their, those people…

Pronoun type ■ Posts

■■■■■■■■■■■■■■■■■■■■■■■■■

none ■ 85 292 (63.3 %)

them only ■ 24 916 (18.5 %)

we only ■ 15 261 (11.3 %)

both ■ 8 990 (6.7 %)

■ **Toxicity scoring** — unitary/toxic-bert (10 k sample)

→ 5 dimensions: toxicity, severe_toxicity,

identity_attack, insult, threat

→ Average toxicity score: **0.593**

■ **Emotion classification** — GoEmotions (10 k sample)

→ top-1 label among 27 emotion classes

→ neutral 35.4 % · anger 18.0 % · annoyance 10.4 %

→ curiosity 5.8 % · admiration 5.0 %

Step 3 — Rule-Based Othering Detector

Each post is scanned against **32 compiled regex patterns** grouped in 4 linguistic families. Each matched family adds 1 point to an **othering_score** (0–4).

Family	Example patterns	Posts matched

Dehumanising metaphors	invasion · flood · swarm · vermin · parasite · animals	5 463 (4.1 %)
Moral exclusion	don't belong · go back · send them back · get out	626 (0.5 %)
Generalisation	all of them · those people · they never · none of them	976 (0.7 %)
Threat framing	taking over · replacing us · great replacement	50 (0.04%)

Top matched patterns across the full dataset:

animals	3,326
invasion	1,104
vermin	1,041
none_of_them	832
parasite	831
go_back	274
all_of_them	150

■ **Othering score distribution**

→ score 0 (no othering) : 127 344 posts (94.7 %)

→ score 1 (1 family) : 6 263 posts (4.7 %)

→ score 2 (2 families) : 849 posts (0.6 %)

→ score 3 (3 families) : 3 posts (< 0.1 %)

■ **Toxicity correlation**

→ Avg toxicity — othering posts : **0.760**

→ Avg toxicity — non-othering posts : 0.583

→ Othering posts are significantly more toxic

■ **By source**: hate_speech 5.3 % · reddit 0.7 %

Next Steps

Step 4	ML classifier — in progress (TF-IDF & sentence embeddings × LR / LinearSVC)
Step 5	Topic modelling (BERTopic) on the 7 115 othering posts → identify recurring narratives
Step 6	Temporal & community analysis — othering rate over time, by subreddit
Step 7	Fine-tuned BERT classifier with a small set of manual annotations